

AI Engineer:

Job Overview: Our team of students works with machine learning, deep learning, and signal analysis for audio, video, and textual analysis in a medical setting. We welcome students with experience in these domains and mentors with deep expertise that can guide us to build a robust solution as soon as possible.

Responsibilities:

- Set up and manage our AI development and production infrastructure.
- Help AI product managers and business stakeholders understand the potential and limitations of AI when planning new products.
- Build data ingest and data transformation infrastructure.
- Identify transfer learning opportunities and new training datasets.
- Build AI models from scratch and help product managers and stakeholders understand results.
- Deploy AI models into production.

Requirements:

- Hands-on programming experience working on enterprise products,
- Demonstrated proficiency in multiple programming languages with a strong foundation in a statistical platform such as Python, R,
- Experience building AI models in platforms such as Keras, TensorFlow, or Theano, experience with computer vision, NLP, signal processing
- Demonstrated commitment to learning about AI through your own initiatives through courses, books, or side project

Hours:

15+ hours per week

To Apply:

Please fill out the [form](#) and feel free to reach out to azehfuss@stanford.edu with any questions